



Phase 2 provider comparison (descriptive, n=300 per language)

Frozen Phase 2 corpus: 300 Uzbek + 300 Kazakh utterances · both providers × both conditions · no inferential tests conducted in Phase 2

A · Mean WER by provider, language, and condition



B · Request latency (Phase 2)



ElevenLabs Scribe v2 outperforms Google Cloud STT Chirp 2 in all four cells and is ~2.6× faster. Phase 2 was descriptive (no pre-specified hypothesis tests); cross-provider WER comparisons here are observational.