


Central Asian Voice Benchmark — Overview

Speech-to-text for Uzbek & Kazakh · frozen results, 2026-08-25 · benchmark only, no product recommendation



What each phase covered

- Phase 1**
3 models · gpt-4o / mini / Whisper-1
context only — different corpus [report]
- Phase 2**
2 providers · ElevenLabs + GCS Chirp 2
600 utts (300 uz + 300 kk) × 2 = 2,400 recs
- Phase 3A · analysis**
ElevenLabs only · detection study
160 eval utts (70 uz + 90 kk)
- Phase 3A · test**
R1 routing vs always-HINT (PRIMARY)
PRIMARY criterion NOT MET (p = 0.820)

Key finding

- Automatic language detection fails for Uzbek (27%) but not Kazakh (99%).
- Misdetected Uzbek transcribes with ~2× the error rate.
- A detect-and-route rule (R1) matched but did not beat always supplying the hint (PRIMARY p = 0.820).