



Model & provider comparison (descriptive)

Lower WER is better · HINT = language supplied, AUTO = auto-detect · Phase 2 = pre-specified descriptive comparison; Phase 3A = single-provider follow-up

A · Phase 2 — two providers, 4 cells each (n = 300)



B · Phase 3A — ElevenLabs Scribe v2 only (n = 70 uz / 90 kk)



PRIMARY vs SECONDARY: Panel B's ROUTED vs HINT is the PRE-SPECIFIED PRIMARY test (null result, p=0.820); AUTO vs HINT in either panel is a SECONDARY description. Phase 1 models (gpt-4o, mini, Whisper-1) use a different corpus and are omitted here — see the report Appendix A for context.
Source: Central Asian Voice Benchmark — frozen Phase 1/2/3A results (final_benchmark_report.md, 2026-08-25) · central-asian-voice-benchmark