


Accuracy vs cost per call

Phase 2 (2 providers × 2 languages × 2 conditions) and Phase 3A (ElevenLabs only). Cost = mean measured per-call USD from frozen JSONL. Bubble size = mean latency. Descriptive.



ElevenLabs Scribe v2 dominates the cost-accuracy frontier: lowest WER in all four Phase 2 cells at ~1/15 the per-call cost of GCS Chirp 2 and 2.6× lower mean latency. Phase 2 and Phase 3A use disjoint utterance sets (same corpus), so cross-phase points are indicative, not directly comparable. Phase 2 was descriptive (no tests).
Source: Central Asian Voice Benchmark — frozen Phase 1/2/3A results (final_benchmark_report.md, 2026-08-25) · central-asian-voice-benchmark