



Latency vs WER (per-utterance)

Phase 3A evaluation (320 records, ElevenLabs Scribe v2) per-utterance · Phase 2 provider means overlaid as labelled points · latency is dominated by audio duration

A · Phase 3A per-utterance (n = 320)



B · Phase 2 provider latency (mean, s)



Phase 3A per-utterance latency varied more by language (Uzbek shorter audio → lower latency) than by condition. HINT latency is marginally lower than AUTO (mean 1.03 s vs 1.07 s). Phase 2 means are provider-level (shared corpus).

Source: Central Asian Voice Benchmark — frozen Phase 1/2/3A results (final_benchmark_report.md, 2026-08-25) · central-asian-voice-benchmark