


Reproducibility & data integrity

All counts recomputed from frozen JSONL at generation time · verification: verify_against_report.py (41/41) + verify_publication.py

A · Dataset composition (frozen JSONL)



B · Integrity checks

(frozen JSONL, no edits)

VERIFY publication	PASS
VERIFY report (41/41)	PASS
Actual spend = \$0.1497	PASS
Uzbek eval = 70 / Kazakh = 90	PASS
Eval AUTO = 160 / HINT = 160	PASS
Pilot = 80 / Eval = 320	PASS
Dup successful tuples = 0	PASS
Historical errors = 68	PASS
Successful = 400	PASS
Total records = 468	PASS

Provenance ledger

Phase 2 version
phase2_trackb_v1

Phase 3A version
phase3a_trackb_v1

Phase 3A frozen commit
b8b364c

Manifest P2 SHA-256
d01408af2e04d048...

Manifest P3A SHA-256
4e2ab89f7c30afeb...

P2 JSONL SHA-256
6a0415cc931352e9...

P3A JSONL SHA-256
8c1cfa6e9b9c03bb...

Full hashes in final_benchmark_report.md §4, §19

Frozen results are treated as immutable research artifacts: this analysis layer makes no API calls, spends no money, and never modifies the source JSONL. 0 duplicate successful tuples; 0 utterance overlap between Phase 2 and Phase 3A. ElevenLabs logs transcripts server-side by default — disclosed per report §19.

Source: Central Asian Voice Benchmark — frozen Phase 1/2/3A results (final_benchmark_report.md, 2026-08-25) · central-asian-voice-benchmark